A simple method for automated equilibration detection in molecular simulations

John D. Chodera[1,*]

[1]*Computational Biology Program, Sloan Kettering Institute, Memorial Sloan Kettering Cancer Center, New York, NY 10065*

(Dated: June 1, 2015)

Molecular simulations intended to compute equilibrium properties are often initiated from configurations that are highly atypical of equilibrium samples, a practice which can generate a distinct initial transient in mechanical observables computed from the simulation trajectory. Traditional practice in simulation data analysis recommends this initial portion be discarded to *equilibration*, but no simple, general, and automated procedure for this process exists. Here, we suggest such a conceptually simple automated procedure that does not make strict assumptions about the distribution of the observable of interest, where the equilibration region is chosen to maximize the number of effectively uncorrelated samples in the production portion used to compute equilibrium averages. We present a simple Python reference implementation of this procedure, and demonstrate its utility on both synthetic and real simulation data.

Keywords: molecular dynamics (MD); Metropolis-Hastings; Monte Carlo (MC); Markov chain Monte Carlo (MCMC); equilibration; burn-in; timeseries analysis; statistical inefficiency; integrated autocorrelation time

INTRODUCTION

Molecular simulations use Markov chain Monte Carlo (MCMC) techniques [1] to sample configurations x from an equilibrium distribution $\pi(x)$, either exactly (using Monte Carlo methods such as Metropolis-Hastings) or approximately (using molecular dynamics integrators without Metropolization) [2].

Due to the sensitivity of the equilibrium density $\pi(x)$ to small perturbations in configuration x and the difficulty of producing sufficiently good guesses of typical equilibrium configurations $x \sim \pi(x)$, these molecular simulations are often started from highly atypical initial conditions. For example, simulations of biopolymers might be initiated from a fully extended conformation unrepresentative of behavior in solution, or a geometry derived from a fit to diffraction data collected from a cryocooled crystal; solvated systems may be prepared by periodically replicating a small solvent box equilibrated under different conditions, yielding atypical densities and solvent structure; liquid mixtures or lipid bilayers may be constructed by using methods that fulfill spatial constraints (e.g. PackMol [3]) but create locally aytpical geometries, requiring long simulation times to relax to typical configurations.

As a result, traditional practice in molecular simulation has recommended some initial portion of the trajectory be discarded to *equilibration* (also called *burn-in*[1] in the MCMC literature [4]). While the process of discarding initial samples is strictly unnecessary for the time-average of quantities of interest to eventually converge to the desired expectations [5], and as a result is often not recommended by statisticians [4], the differences in complexity of probability densities typically encountered in statistics and molecular simulation may explain the difference in historical practice. In the simulation field, discarding initial samples to equilibration nevertheless often allows the practitioner to avoid what may be impractically long run times to eliminate the bias in computed properties in finite-length simulations induced by atypical initial starting conditions.

As an illustrative example, consider the computation of the average density of liquid argon under a given set of reduced temperature and pressure conditions Figure 1. To initiate the simulation, an initial dense liquid geometry at reduced density $\rho^* \equiv \rho\sigma^3 = 0.960$ was prepared and subjected to local energy minimization. Figure 1 (top) depicts the relaxation behavior of 100 simulations initiated from the same configuration with different random initial velocities and integrator random number seeds (full simulation details are provided in *Simulation Details*). The average (black line) and standard deviation (shaded grey) shows that all realizations of this simulation show a characteristic relaxation behavior away from the initial density toward the equilibrium density. The expectation of the running average of the density over many realizations of this procedure (Figure 1, bottom) significantly deviates from the actual expectation, which would lead to biased estimates unless simulations were sufficiently long to eliminate this starting point dependent bias. Note that this significant bias is present because the *same* atypical starting condition is used for every realization of this simulation process.

STATEMENT OF THE PROBLEM

Consider successively sampled configurations x_t from a molecular simulation, with $t = 1, \ldots, T$. We presume we are interested in computing the expectation $\langle A \rangle \equiv \int dx\, A(x)\, \pi(x)$ of a mechanical property $A(x)$. For convenience, we will refer to the timeseries $a_t \equiv A(x_t)$, with $t = 0$. The estimator $\hat{A} \approx \langle A \rangle$ constructed from the entire

* Corresponding author; john.chodera@choderalab.org

[1] The term *burn-in* comes from the field of electronics, in which a short "burn-in" period is used to ensure that a device is free of faulty components—which often fail quickly—and is operating normally [4].

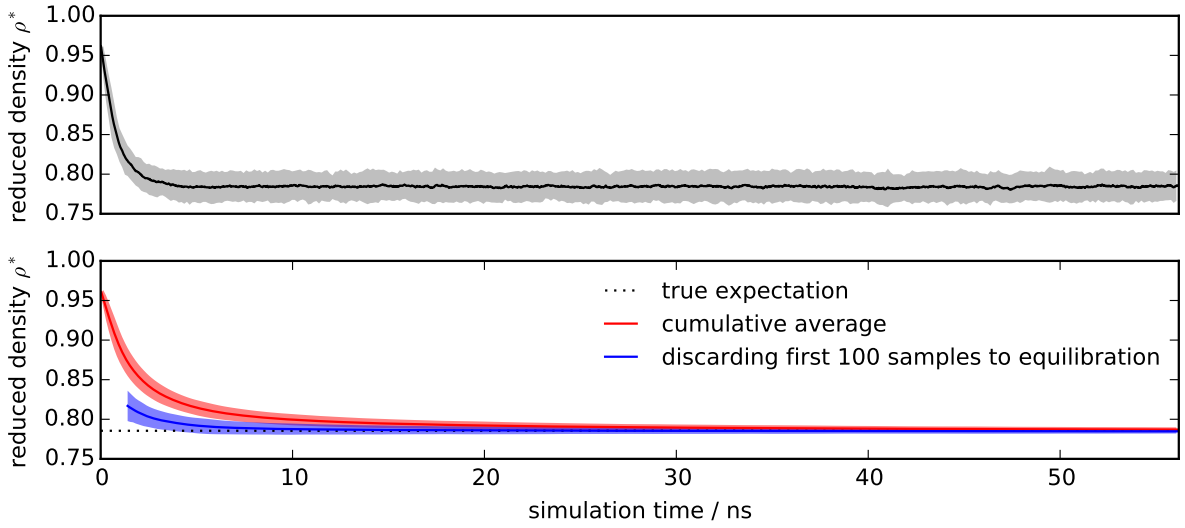

FIG. 1. Illustration of the motivation for discarding data to equilibration. To illustrate the bias in expectations induced by relaxation away from initial conditions, 100 replicates of a simulation of liquid argon were initiated from the same energy-minimized initial configuration constructed with initial reduced density $\rho^* \equiv \rho\sigma^3 = 0.960$ but different random number seeds for stochastic integration. **Top:** The average of the reduced density (black line) over the replicates relaxes to the region of typical equilibrium densities over the first few ns of simulation time. **Bottom:** If the average density is estimated by a cumulative average from the beginning of the simulation (red line), the estimate will be heavily biased by the atypical starting density even beyond 10 ns. Discarding even a small amount of initial data—in this case 100 initial samples (blue line)—results in a cumulative average estimate that converges to the true average (black dotted line) much more rapidly. Shaded regions denote 95% confidence intervals.

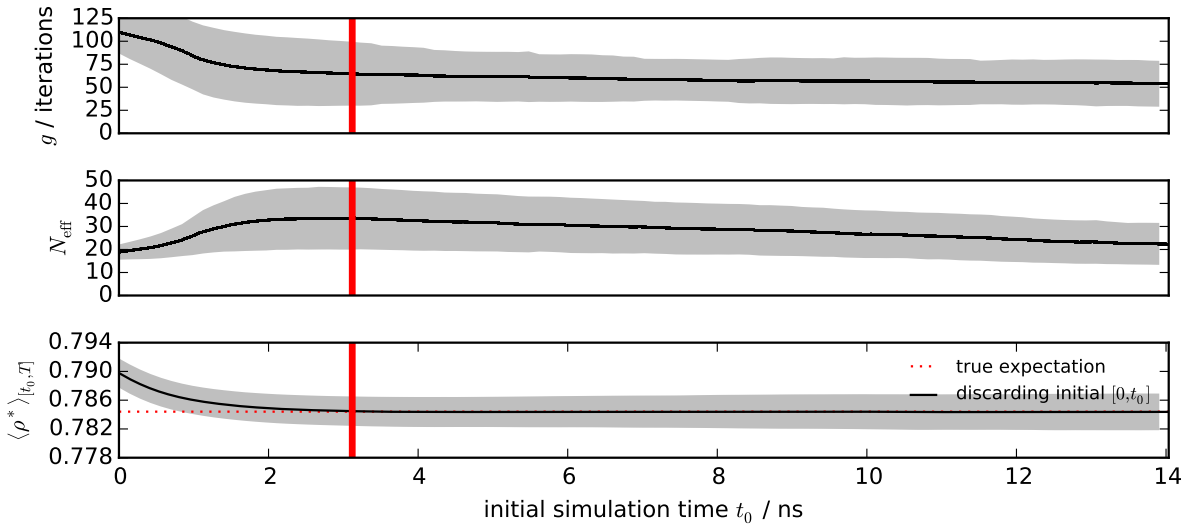

FIG. 2. Statistical inefficiency, number of uncorrelated samples, and bias for different burn-in times. Trajectories of length $T = 2\,000$ iterations (∼28 ns) for the Lennard-Jones system described in Fig. 1 were analyzed as a function of equilibration time choice t_0. Averages over all 100 replicate simulations (all starting from the same initial conditions) are shown as dark lines, with shaded lines showing standard deviation of estimates among replicates. **Top:** The statistical inefficiency g as a function of equilibration time choice t_0 is initially very large, but diminishes rapidly after the system has relaxed to equilibrium. **Middle:** The number of effectively uncorrelated samples $N_{\mathrm{eff}} = (T - t_0 + 1)/g$ shows a maximum at $t_0 = 222$ iterations, suggesting the system has equilibrated by this time. The red vertical line in all plots marks this choice of $t_0 = 222$. **Bottom:** The cumulative density average $\langle\rho^*\rangle$ computed over the span $[t_0, T]$ shows that the bias (deviation from the true estimate, shown as red dashed lines) is minimized for choices of $t_0 \geq 222$ iterations. The standard deviation among replicates (shaded region) grows with t_0 because fewer data are included in the estimate. The choice of optimal t_0 that maximizes N_{eff} (red vertical line) strikes a good balance between bias and variance. The true estimate (red dashed lines) is computed from averaging over the range [5 000, 10 000] iterations over all 100 replicates.

dataset is given by

$$\hat{A}_{[1,T]} \equiv \frac{1}{T}\sum_{t=1}^{T} a_t. \tag{1}$$

While $\lim_{T\to\infty} \hat{A}_{[1,T]} = \langle A \rangle$ for an infinitely long simulation[2], the bias in $\hat{A}_{[1,T]}$ may be significant in a simulation of finite length T.

By discarding samples $t < t_0$ to equilibration, we hope to eliminate the initial transient and provide a less biased estimate of $\langle A \rangle$,

$$\hat{A}_{[t_0,T]} \equiv \frac{1}{T-t_0+1}\sum_{t=t_0}^{T} a_t. \tag{2}$$

We can quantify the bias in an estimator $\hat{A}$ by the expected error $\delta^2\hat{A}$,

$$\delta^2\hat{A} \equiv E_{x_0}\left[\left(\hat{A} - \langle A \rangle\right)^2\right]. \tag{3}$$

where $E_{x_0}[\cdot]$ denotes the expectation over independent realizations of the simulation from the same initial configuration x_0 but different initial velocities and random number seeds.

Here, we concern ourselves with this question: Is there a simple approach to choosing an optimal equilibration time t_0 that provides an improved estimate $\hat{A}_{[t_0,T]}$, such that $\delta^2\hat{A}_{[t_0,T]} < \delta^2\hat{A}_{[1,T]}$? We note that, for cases in which the simulation is not long enough to reach equilibrium, no choice of t_0 will eliminate bias completely; the best we can hope for is to minimize this bias.

While several automated methods for selecting the equilibration time t_0 have been proposed, these approaches have shortcomings that have greatly limited their use. The reverse cumulative averaging method [6], for example, uses a statistical test for normality to determine the point before which which the observable timeseries deviates from normality. While this concept may be reasonable for experimental data, where measurements often represent the sum of many random variables such that the central limit theorem's guarantee of asymptotic normality ensures the distribution of the observable will be approximately normal, there is no such guarantee that instantaneous measurements of a simulation property of interest will be normally distributed. In fact, many properties will be decidedly *non-normal*. For a biomolecule such as a protein, for example, the radius of gyration, end-to-end distance, and torsion angles sampled during a simulation will all be highly non-normal. Instead, we require a method that makes no assumptions about the nature of the distribution of the property under study.

[2] We note that this equality only holds for simulation schemes that sample from the true equilibrium distribution $\pi(x)$, such as Metropolis-Hastings Monte Carlo or Metropolized integration schemes. Molecular dynamics simulations utilizing finite timestep integration without Metropolization will produce averages that may deviate from the true expectation $\langle A \rangle$ [**?**].

AUTOCORRELATION ANALYSIS

The set of successively sampled configurations $\{x_t\}$ and their corresponding observables $\{a_t\}$ compose a correlated timeseries of observations. To estimate the statistical error or uncertainty in a stationary timeseries free of bias, we must be able to quantify the *effective number of uncorrelated samples* present in the dataset. This is usually accomplished through computation of the *statistical inefficiency* g, which quantifies the number of correlated timeseries samples needed to produce a single effectively uncorrelated sample of the observable of interest. While these concepts are well-established for the analysis of both Monte Carlo and molecular dynamics simulations [7–10], we review it here for the sake of clarity.

For a given equilibration time choice t_0, the statistical uncertainty in our estimator $\hat{A}_{[t_0,T]}$ can be written as,

$$\begin{aligned}
\delta^2\hat{A}_{[t_0,T]} &\equiv E_{x_0}\left[\left(\hat{A}_{[t_0,T]} - \langle\hat{A}\rangle\right)^2\right] \\
&= E_{x_0}\left[\hat{A}^2_{[t_0,T]}\right] - E_{x_0}\left[\hat{A}_{[t_0,T]}\right]^2 \\
&= \frac{1}{T^2_{t_0}}\sum_{t,t'=t_0}^{T}\left[\langle a_t a_{t'}\rangle - \langle a_t\rangle\langle a_{t'}\rangle\right] \\
&= \frac{1}{T^2_{t_0}}\sum_{t=t_0}^{T}\left[\langle x_t^2\rangle - \langle x_t\rangle^2\right] \\
&\quad + \frac{1}{T^2_{t_0}}\sum_{t\neq t'=t_0}^{T}\left[\langle a_t a_{t'}\rangle - \langle a_t\rangle\langle a_{t'}\rangle\right].
\end{aligned} \tag{4}$$

where $T_{t_0} \equiv T - t_0 + 1$, the number of correlated samples in the timeseries $\{a_t\}_{t_0}^{T}$. In the last step, we have split the double-sum into two separate sums—a term capturing the variance in the observations a_t, and a remaining term capturing the correlation between observations.

If t_0 is sufficiently large for the initial bias to be eliminated, the remaining timeseries $\{a_t\}_{t_0}^{T}$ will obey the properties of both *stationarity* and *time-reversibility*, allowing us to write,

$$\begin{aligned}
\delta^2\hat{A}^{\text{equil}}_{[t_0,T]} &= \frac{1}{T_{t_0}}\left[\langle a_t^2\rangle - \langle a_t\rangle^2\right] \\
&+ \frac{2}{T_{t_0}}\sum_{n=1}^{T-t_0}\left(\frac{T_{t_0}-n}{T_{t_0}}\right)\left[\langle a_t a_{t+n}\rangle - \langle a_t\rangle\langle a_{t+n}\rangle\right] \\
&\equiv \frac{\sigma^2_{t_0}}{T_{t_0}}(1+2\tau_{t_0}) \\
&= \frac{\sigma^2_{t_0}}{T_{t_0}/g_{t_0}}
\end{aligned} \tag{5}$$

where the variance σ^2, statistical inefficiency g, and integrated autocorrelation time τ (in units of the sampling in-

terval) are given by

$$\sigma^2 \equiv \langle a_t^2 \rangle - \langle a_t \rangle^2 \tag{6}$$

$$\tau \equiv \sum_{t=1}^{T-1} \left(1 - \frac{t}{T}\right) C_t \tag{7}$$

$$g \equiv 1 + 2\tau \tag{8}$$

with the discrete-time normalized fluctuation autocorrelation function C_t defined as

$$C_t \equiv \frac{\langle a_n a_{n+t} \rangle - \langle a_n \rangle^2}{\langle a_n^2 \rangle - \langle a_n \rangle^2}. \tag{9}$$

In practice, it is difficult to estimate C_t for $t \sim T$, due to growth in the statistical error, so common estimators of g make use of several additional properties of C_t to provide useful estimates [4].

The t_0 subscript for the variance σ^2, the integrated autocorrelation time τ, and the statistical inefficiency t_0 mean that these quantities are only estimated over the production portion of the timeseries, $\{a_t\}_{t=t_0}^{T}$. Since we assumed that the bias was eliminated by judicious choice of the equilibration time t_0, this estimate of the statistical error will be poor for choices of t_0 that are too small.

THE ESSENTIAL IDEA

Suppose we choose some arbitrary time t_0 and discard all samples $t \in [0, t_0)$ to equilibration, keeping $[t_0, T]$ as the dataset to analyze. How much data remains? We can determine this by computing the statistical inefficiency g_{t_0} for the interval $[t_0, T]$, and computing the effective number of uncorrelated samples $N_{\mathrm{eff}}(t_0) \equiv (T - t_0 + 1)/g_{t_0}$. If we start at $t_0 \equiv T$ and move t_0 to earlier and earlier points in time, we expect that the effective number of uncorrelated samples $N_{\mathrm{eff}}(t_0)$ will continue to grow until we start to include the highly atypical initial data. At that point, the integrated autocorrelation time τ (and hence the statistical inefficiency g) will greatly increase, and the effective number of samples N_{eff} will start to plummet.

This suggests an alluringly simple algorithm for identifying the optimal equilibration time—pick the t_0 which maximizes the number of uncorrelated samples N_{eff}. In mathematical terms,

$$t_0^{\mathrm{opt}} = \operatorname*{argmax}_{t_0} N_{\mathrm{eff}}(t_0) \tag{10}$$

$$= \operatorname*{argmax}_{t_0} \frac{T - t_0 + 1}{g_{t_0}} \tag{11}$$

Figure 2 demonstrates this for the liquid argon system described above, using expectations computed over 100 independent replicate trajectories. At short t_0, the statistical inefficiency g (Figure 2, top panel) is large due to the contribution from slow relaxation from atypical initial conditions, while at long t_0 the statistical inefficiency estimate is much shorter and nearly constant of a large span of time origins. As a result, the effective number of uncorrelated samples N_{eff} (Figure 2, middle panel) has a peak at $t_0 \sim 222$ iterations (Figure 2, vertical red lines). The effect on bias in the estimated average reduced density $\langle \rho^* \rangle$ (Figure 2, bottom panel) is striking—the bias is essentially eliminated for the choice of equilibration time t_0 that maximizes the number of uncorrelated samples N_{eff}.

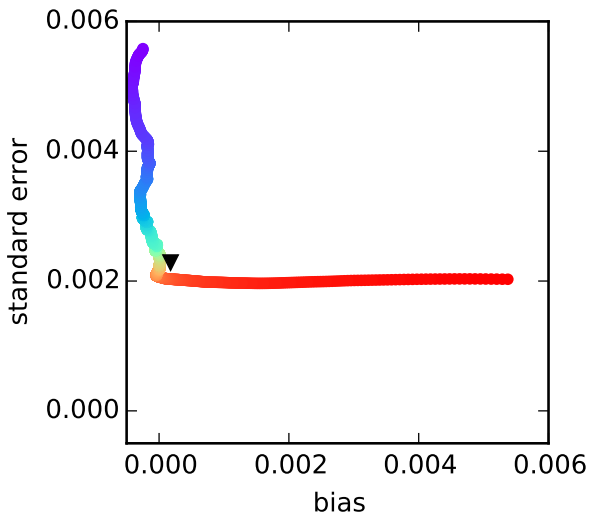

FIG. 3. Bias-variance tradeoff for fixed equilibration time versus automatic equilibration time selection. Trajectories of length T = 2 000 iterations (∼28 ns) for the Lennard-Jones system described in Fig. 1 were analyzed as a function of equilibration time choice t_0. Using 100 replicate simulations, the average bias (average deviation from true expectation) and standard deviation (random variation from replicate to replicate) were computed as a function of a prespecified fixed equilibration time t_0, with colors running from $t_0 = 0$ (red) to t_0 =1 800 iterations (blue). As is readily discerned, the bias for small t_0 is initially large, but minimized for larger t_0. By contrast, the standard error (a measure of variance, estimated here by standard deviation among replicates) grows as t_0 grows above a certain critical time (here, ∼222 iterations). If the t_0 that maximizes N_{eff} is instead chosen *individually* for each trajectory based on that trajectory's estimated statistical inefficiency $g_{[t_0,T]}$, the resulting bias-variance tradeoff (black triangle) does an excellent job minimizing bias and variance simultaneously, comparable to what is possible for a choice of equilibration time t_0 based on knowledge of the true bias and variance among many replicate estimates. [JDC: Add colorbar.]

BIAS-VARIANCE TRADEOFF

With increasing equilibration time t_0, bias is reduced, but the variance—the contribution to error due to random variation from having a finite number of uncorrelated samples—will increase because less data is included in the estimate. This can be seen in the bottom panel of Figure 2, where the shaded region (denoting the standard deviation among

sample estimates) increases in width with increasing equilibration time t_0.

To examine the tradeoff between bias and variance explicitly, Figure 4 plots the bias and variance (here, shown as standard error) contributions against each other as a function of t_0 (denoted by color). At $t_0 = 0$, the bias is large but variance is minimized. With increasing t_0, bias is eventually eliminated but then variance rapidly grows as fewer uncorrelated samples are included in the estimate. There is a clear optimal choice at $t_0 \sim 222$ iterations that minimizes variance while also effectively eliminating bias.

But how will this strategy work for cases where we do not know the statistical inefficiency g as a function of the equilibration time t_0 precisely? When all that is available is a single simulation, our best estimate of g_{t_0} is estimated from that simulation alone over the span $[t_0, T]$—will this affect the quality of our estimate of equilibration time? Empirically, this does not appear to be the case—the black triangle in Figure 4 shows the bias and variance for estimates computed over the 100 replicates where t_0 is individually determined from each simulation.

DISCUSSION

The scheme described here—in which the equilibration time t_0 is computed using Eq. 10 as the time origin that maximizes the number of uncorrelated samples in the production region $[t_0, T]$—is both conceptually and computationally straightforward. It provides an approach to determining the optimal amount of initial data to discard to equilibration in order to minimize variance while also minimizing initial bias, and does this without employing statistical tests that require unsatisfiable assumptions of normality of the observable of interest. As we have seen, this scheme empirically appears to select a practical compromise between bias and variance even when the statistical inefficiency g is estimated directly from the trajectory using Eq. 8.

A word of caution is necessary. One can certainly envision pathological scenarios where this algorithm for selecting an optimal equilibration time will break down. In cases where the simulation is not long enough to reach equilibrium—let alone collect many uncorrelated samples from it—no choice of equilibration time will bestow upon the data the ability to produce an unbiased estimate of the true expectation. Similarly, in cases where insufficient data is available for the statistical inefficiency to be estimated well, this algorithm is expected to perform poorly. However, in these cases, the data itself should be suspect if the trajectory is not at least an order of magnitude longer than the minimum estimated autocorrelation time.

SIMULATION DETAILS

All molecular dynamics simulations described here were performed with OpenMM 6.2 [11] (available at openmm.org) using the Python API. All scripts used to run simulations, analyze data, and generate plots—along with the simulation data itself and scripts for generating figures—are availabile on GitHub[3].

The argon model system comes from the openmmtools package[4]. Simulations were performed using a box of $N = 500$ argon atoms at reduced temperature $T^* \equiv k_B T/\epsilon = 0.850$ and reduced pressure $p^* \equiv p\sigma^3/\epsilon = 1.266$ using a Langevin integrator [12] with timestep $\Delta t = 0.01\tau$, where characteristic oscillation timescale $\tau = \sqrt{mr_0^2/72\epsilon}$, with $r_0 = 2^{1/6}\sigma$ [13]. A Metropolis Monte Carlo barostat was used with box volume moves attempted every 25 timesteps. Densities were recorded every 25 timesteps.

The automated equilibration detection scheme is also available in the `timeseries` module of the `pymbar` package as `detectEquilibration()`, and can be accessed using the following code:

```
from pymbar.timeseries import detectEquilibration
# determine equilibrated region
[t0, g, Neff_max] = detectEquilibration(A_t)
# discard initial samples to equilibration
A_t = A_t[t0:]
```

PRACTICAL COMPUTATION OF STATISTICAL INEFFICIENCY

The computation of the statistical inefficiency g (defined by Eq. 8) for a finite timeseries a_t, $t = 1, \ldots, T$ deserves some comment. There are, in fact, a variety of schemes for estimating g described in the literature, and their behaviors for finite datasets may differ, leading to different estimates of the equilibration time t_0 using the algorithm of Eq. 10.

The main issue is that a straightforward approach to estimating the statistical inefficiency using Eqs. 7–9 in which the expectations are simply replaced with sample estimates causes the statistical error in the estimated correlation function C_t to grow with t in a manner that allows this error to quickly overwhelm the sum of Eq. 7. As a result, a number of alternative schemes—generally based on controlling the error in the estimated C_t or truncating the sum of Eq. 7 when the error grows too large—have been proposed.

For stationary, reversible Markov chains, Geyer observed that a function $\Gamma_k \equiv \gamma_{2k} + \gamma_{2k+1}$ of the unnormalized fluctuation autocorrelation function $\gamma_t \equiv \langle a_i a_{i+t} \rangle - \langle a_i \rangle^2$ has a number of pleasant properties (Theorem 3.1 of [14]): It is strictly, positive, strictly decreasing, and strictly convex. [JDC: Check conditions of Geyer proofs.] These properties can be exploited to define a family of estimators called *initial sequence methods* (see Section 3.3 of [14] and Section 1.10.2 of [4]), of which the *initial convex sequence* (ICS) estimator

[3] All scripts and data are available at: http://github.com/choderalab/automatic-equilibration-detection
[4] available at http://github.com/choderalab/openmmtools

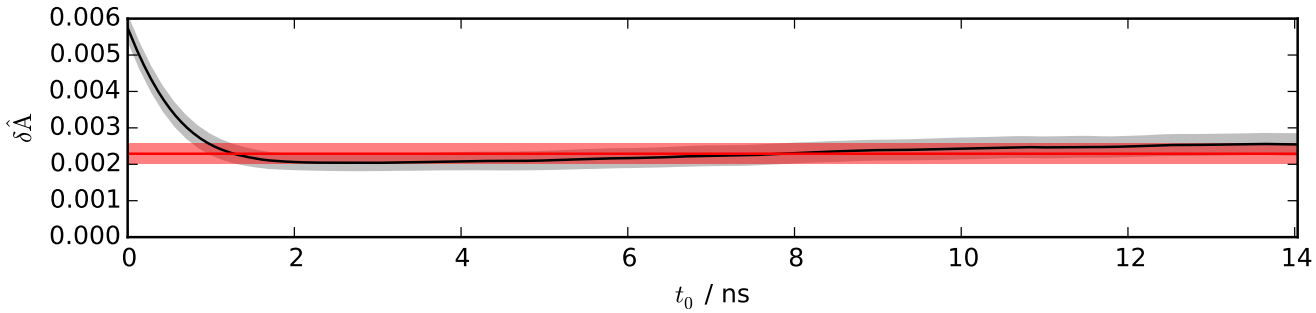

FIG. 4. RMS error for fixed equilibration time versus automatic equilibration time selection. Trajectories of length $T = 2\,000$ iterations (∼28 ns) for the Lennard-Jones system described in Fig. 1 were analyzed as a function of fixed equilibration time choice t_0. Using 100 replicate simulations, the RMS error (average root mean squared deviation from true expectation, as defined in Eq. 3) was computed (black line) along with 95% confidence interval (gray shading). The RMS error is minimized for fixed equilibration time choices in the range 2–6 ns. If the t_0 that maximizes N_{eff} is instead chosen *individually* for each trajectory based on that trajectory's estimated statistical inefficiency $g_{[t_0,T]}$ using Eq. 10, the resulting RMS error (red line, 95% confidence interval shown as red shading) is quite close to the minimum RMS error achieved from any particular fixed equilibration time t_0, suggesting that this simple automated approach to selecting t_0 performs reasonably well.

is generally agreed to be optimal, if somewhat complex to implement. [JDC: Give implementation details?]

All computations in this manuscript used the fast multiscale method described in Section 5.2 of [10]. This method is related to a multiscale variant of the *initial positive sequence* (IPS) method of Geyer [15], where contributions are accumulated at increasingly longer lag times and the sum of Eq. 7 is truncated when the terms become negative. We have found this method to be both fast and to provide useful estimates of the statistical inefficiency, but it may not perform well for all problems. [JDC: Give implementation details? Compare with ICS?]

ACKNOWLEDGMENTS

We are grateful to William C. Swope (IBM Almaden Research Center), Michael R. Shirts (University of Virginia), David L. Mobley (University of California, Irvine), Kyle A. Beauchamp (MSKCC), and Robert C. McGibbon (Stanford University) for valuable discussions on this topic, and Joshua L. Adelman (University of Pittsburgh) for helpful feedback and encouragement.

[1] J. S. Liu, *Monte Carlo strategies in scientific computing*, 2nd ed. ed. (Springer-Verlag, New York, 2002).

[2] D. Sivak, J. Chodera, and G. Crooks, Physical Review X **3**, 011007 (2013), bibtex: Sivak:2013:Phys.Rev.X.

[3] L. Martínez, R. Andrade, E. G. Birgin, and J. M. Martínez, J. Chem. Theor. Comput. **30**, 2157 (2009).

[4] S. Brooks, A. Gelman, G. L. Jones, and X.-L. Meng, in *Handbook of Markov chain Monte Carlo*, *Chapman & Hall/CRC Handbooks of Modern Statistical Methods* (CRC Press, ADDRESS, 2011), Chap. Introduction to Markov chain Monte Carlo.

[5] C. Geyer, Burn-in is unnecessary., http://users.stat.umn.edu/~geyer/mcmc/burn.html.

[6] W. Yang, R. Bittetti-Putzer, and M. Karplus, J. Chem. Phys. **120**, 2618 (2004).

[7] H. Müller-Krumbhaar and K. Binder, J. Stat. Phys. **8**, 1 (1973).

[8] W. C. Swope, H. C. Andersen, P. H. Berens, and K. R. Wilson, J. Chem. Phys. **76**, 637 (1982).

[9] W. Janke, in *Quantum Simulations of Complex Many-Body Systems: From Theory to Algorithms*, edited by J. Grotendorst, D. Marx, and A. Murmatsu (John von Neumann Institute for Computing, ADDRESS, 2002), Vol. 10, pp. 423–445.

[10] J. D. Chodera, W. C. Swope, J. W. Pitera, C. Seok, and K. A. Dill, J. Chem. Theor. Comput. **3**, 26 (2007).

[11] P. Eastman, M. Friedrichs, J. D. Chodera, R. Radmer, C. Bruns, J. Ku, K. Beauchamp, T. J. Lane, L.-P. Wang, D. Shukla, T. Tye, M. Houston, T. Stitch, and C. Klein, J. Chem. Theor. Comput. **9**, 461 (2012).

[12] D. A. Sivak, J. D. Chodera, and G. E. Crooks, J. Phys. Chem. B **118**, 6466 (2014).

[13] B. Veytsman and M. Kotelyanskii, Lennard-Jones potential revisited., http://borisv.lk.net/matsc597c-1997/simulations/Lecture5/node3.html.

[14] C. J. Geyer, Stat. Sci. **76**, 473 (1992).

[15] C. J. Geyer and E. A. Thompson, J. Royal Stat. Soc. B **54**, 657 (1992).